


```
In [62]: print("Number of neurons in the inner layer: ", node_size)
        for (a, l, bn_l) in losses:
            # Visualize
            plt.figure(figsize=(15, 10))

            plt.title("Training loss")
            plt.xlabel("#iteration")
            plt.ylabel("loss")
            line = plt.plot(l, 'b', label=a)
            line_bn = plt.plot(bn_l, 'r', label=a + "_BN")

            plt.legend(loc="best")
            plt.show()


            print(a, " loss: ", np.min(l))
            print(a, " loss BatchNorm : ", np.min(bn_l))
```

Number of neurons in the inner layer: 40




ReLU loss: 0.11208550852318845

ReLU loss BatchNorm : 0.10756446750387683




ELU loss: 0.15029989994477483

ELU loss BatchNorm : 0.11749973847113164



LeakyReLU loss: 0.1432750046152047

LeakyReLU loss BatchNorm : 0.09884686500492358



SoftPlus loss: 0.15410253649953057

SoftPlus loss BatchNorm : 0.11186484753411335

In [ ]:

Write your personal opinion on the activation functions, think about computation times too. Does BatchNormalization help?

In [13]: 1 # Your answer goes here. #####

**Finally**, use all your knowledge to build a super cool model on this dataset, do not forget to split dataset into train and validation. Use **dropout** to prevent overfitting, play with **learning rate decay**. You can use **data augmentation** such as rotations, translations to boost your score. Use your knowledge and imagination to train a model. Don't forget to call `training()` and `evaluate()` methods to set desired behaviour of BatchNormalization and Dropout layers.

In [14]: 1 # Your code goes here. #####

Print here your accuracy. It should be around 90%.

In [15]: 1 # Your answer goes here. #####

## Autoencoder